

Cyber5 Command

One skill for the whole Cyber5 season on the Meta Ads MCP. It grades the account, plans the week in your own campaign structure, paces live spend by hand or on a schedule, and reviews the window. It proposes every budget move and makes none.

01 Ready

02 Plan

03 Live

04 Review

THE PROBLEM

The five biggest days break overnight.

Today that means someone watching Ads Manager at 3am and a pacing sheet rebuilt by hand. Last year's Cyber5, replayed on two live accounts:

11h

Cyber Monday ran at near-\$0 delivery while purchases kept coming in.

9pm

Black Friday promo budgets ran dry, three hours before midnight.

32%

of purchases arrived in the evening, on just 24% of spend.

Anonymized external accounts. Each finding came out of the skill's own replay of real accounts.

HOW IT'S ORGANIZED

One skill. Four stages.

It picks the stage from what you ask and today’s date, and names it in the first line of every answer.


01 Ready OCT – EARLY NOV A–F grade across 19 checks, with a dated fix list and an owner for every fix. WRITES TO META Never	02 Plan WEEK OF NOV 16 Last year replayed hour by hour into this year’s plan, in your campaigns and funnel roles. WRITES TO META Never	03 Live NOV 26 – 30 Check-ins by hand, or a morning report on a schedule. Pace, sold-out spend, proposed moves. WRITES TO META Only moves you approve	04 Review ANY PAST WINDOW Replays a window on a simulated clock and scores what it would have caught. WRITES TO META Never
---	--	---	--

Know what breaks before it breaks.

19 read-only checks in five weighted categories. Every failure becomes a fix with an owner and a do-by date, counted back from Thanksgiving. Audiences you don't use aren't failed.


Signal	30	Pixel and server events, match quality
Catalog	20	Feed freshness, must-fix items, pixel match
Delivery	25	Ad errors, disapprovals, billing
Learning	10	Learning stage, recent significant edits
Audiences	15	Warm pools you target, sized for five days

SAMPLE ACCOUNT · 52 DAYS OUT



12 open fixes · first due Nov 12 · each with an owner

Last year, hour by hour, becomes this year's plan.



Purchases by hour, last year's Cyber5 · sample account, anonymized. Outlined: Saturday's delivery gap.

Saturday's budgets ran out at 21:00

Purchases kept coming until midnight. Guard: size Saturday's budgets to last.

Evenings were underfunded

32% of purchases on 24% of spend. Hold 9 points of each day for 18–24.

Split across today's campaigns

Last year sets the shape. This year's money goes to the campaigns running now, ending in a C5PLAN block.

It speaks your naming convention.

YOU SAY

“Our naming: ACQ = acquisition, RTG = remarketing, RET = retention. Targets 60 / 25 / 15.”

IT DOES

Tags every running campaign by role and shows its guesses.

Splits the plan by role and scales each role to your target.

Shows Live pace by role and flags a role drifting 5+ points.

Reroutes money only into acquisition campaigns, by default.

CAMPAIGN	ROLE
ACQ Prospecting Broad	Acquisition
ACQ ASC Evergreen	Acquisition
RTG DPA 30D Visitors	Remarketing
RET Loyalty Existing	Retention
Holiday Gift Guide	Unassigned
No token matched: it says so, and you name the role in one line.	
PLAN BY ROLE	60 / 25 / 15 · matched to your targets

When it lands, and when to type it in.

CAMPAIGN	CHANGE	LANDS AT	ENTER BY
ACQ Prospecting Broad Acquisition	\$4,000 → \$6,500 / day	Fri 00:00	Thu 21:00
RTG DPA 30D Visitors Remarketing	\$1,200 → \$1,800 / day	Fri 00:00	Thu 21:00
ACQ ASC Evergreen Acquisition	+\$900 evening reserve	Sat 18:00	Sat 15:00
RET Loyalty Existing Retention	\$800 → \$1,400 / day may reset learning	Mon 00:00	Sun 21:00

Example plan output, at the 3-hour default lead time.

Lead time is a setting

Enter by = lands at minus budget_lead_hours, 3 by default. Set it to what your team sees in Meta.

Enter the night before

Midnight changes go in the evening before, so the new budget is live when the day starts.

Schedule the five days

Plan points to Meta’s budget scheduling, set by the edit freeze. The skill never sets it for you.

Every check-in runs three passes.

PASS 1

Pace

Projects each campaign's finish from its own hourly shape, and rolls pace up by funnel role. Flags Not started, Stalled and early cap-out first.

PASS 2

Sold-out spend

Maps ads to catalog products and rolls stock up from sizes to products. Names sold-out styles; never counts an unmapped ad as waste.

PASS 3

Reroute

Moves budget from ahead-and-inefficient to behind-and-efficient, and says when to enter each move. Capped at 20% a move, 12-hour cool-down, total stays flat.

THEN

You answer yes or no to each move. Nothing changes until you do.

Ends with a C5LOG block for the next check-in

A baseline before you log on.

MORNING REPORT · FRI NOV 28, 07:00 · FOR THE BUYER

OVERNIGHT

Nothing stalled overnight

YESTERDAY

Thanksgiving: 96% of plan spent, purchases on pace
Overflow finished 18% behind plan, best CPA in its group

TODAY

1 Raise Overflow +\$480/day · enter now
2 Cut Bid-Cap Test -\$480/day · enter now
Open a chat to approve.

Client update draft and C5LOG block attached

Example report and draft, not account data.

CLIENT UPDATE · DRAFT FOR THE BUYER TO FORWARD

Hi Sam, quick Cyber5 update. Thanksgiving spent 96% of plan with purchases on pace (ROAS 4.1, 7-day click). Today we lean into Black Friday as planned. One ask: can you confirm restock timing on the two sold-out hero styles? Next update tomorrow morning.

Same numbers as the report. No budget moves or rankings. Never sent by the skill.

Overnight watch

Hourly on Black Friday and Cyber Monday. Posts only if spending stops.

Read-only, always

Scheduled runs never write, and never message clients.

THE SAFETY MODEL

Propose. Never act.

Built for the five days where one wrong change is most expensive. The skill reads everything and changes nothing on its own.

01

One file can write

Only the approval step can call a Meta write tool. Every other stage is read-only.

02

A yes for every move

No auto-execute, no standing approval. A pasted block, an earlier run or a schedule never counts as a yes.

03

Scheduled runs only report

A morning report lists moves for a person. It can't load the approval step.

04

Re-checked before writing

Each approved move is re-read live first, then logged.

05

Total budget stays flat


Reroutes move money between campaigns. They never raise the Cyber5 total.

06

Never guesses

A failed tool is "couldn't check". An unmatched ad is never called waste. Attribution on every efficiency figure.

It would have caught Cyber Monday's outage at 3am.



Cumulative spend, last year's Cyber5 · sample account, anonymized · plan = each campaign's average day the week before

03:00

Both promo campaigns flagged. All 10 campaigns by 05:00.

8h

before spend came back at 11:00, on the biggest day of the year.

2,531

purchases came in during the dark hours with no ad support.

WHAT TESTING FOUND

Tested on three live accounts, read-only every time.

9pm

Black Friday promo budgets found running dry while purchases kept coming.

16 · 4

real catches against false alarms across the five-day replay.

46% → 4.5%

error in today’s projected finish, 9am to 9pm. So mornings show a range.

0

unplanned Meta MCP calls across every test run, full adherence to skill guidance.

30+

fixes fed back across seven versions, the last from a paid-team review.

1

perfect cold run of the installed skill in a fresh chat, with no prior context.

THREE OUTPUTS PER RUN

One run, three audiences.

Chat view

FOR THE MEDIA BUYER

Headline strip, burn-down and a numbered move list with enter-by times, ending in one approval question.

Stage 03 Live — Nov 27, 11:00

1 Cut Bid-Cap Test -\$480/day · now
2 Raise Evergreen Overflow +\$480/day · now
Approve 1-2? yes / no / pick

Handoff file (.md)

FOR A TEAMMATE OR AN AGENT

Dated tasks with owners, open proposals, and the plan and log blocks, so the next run, a teammate or an agent picks up where this one stopped.

> This file is a record,
> not an authorization.
| T1 | Nov 19 | Backup payment... |
| T6 | Nov 19 | Promo budgets... |

Leadership deck (PDF)

FOR ACCOUNT AND AGENCY LEADS

Up to six slides for Ready, Plan and Review. Same numbers as chat, checked by eye before it's sent. Live stays in chat.

All ten campaigns
flagged by 05:00
Burn-down · fixes for this year
4 slides · 16:9 · PDF

HOW TO USE IT

Ask in plain words. It picks the stage.

- 1

Upload the skill folder in Claude under Customize → Skills.
- 2

Connect the Meta Ads MCP. Nothing else is needed.
- 3

Ask a question. Code execution adds the PDF deck; without it, everything else still works.
- 4

After your first Plan, paste its C5PLAN into a scheduled task for morning reports.

"Is our account ready for Cyber5?"	01 READY
"Build my plan. ACQ = acquisition, RTG = remarketing, RET = retention."	02 PLAN
"Nov 27, 11am. How are we pacing?"	03 LIVE
"Morning report for our account, read-only."	03 LIVE · SCHEDULED
"What went wrong last Cyber5?"	04 REVIEW

Between check-ins: Plan prints a C5PLAN block → paste it into the first Live check-in → every check-in prints a C5LOG block → paste it into the next.

BUILT FOR AGENCIES

Open, tunable, and yours to run.

Meta Ads MCP only

No Shopify, analytics or spreadsheet connections to set up.

Every threshold in one file

Neutral defaults in config/tuned-values.md. Override any in a sentence; changes show up in the caveats.

Your naming, your timing

State your naming convention, role targets and budget lead time in one line.

Sectioned, not monolithic

Each stage loads only the files it needs, in five parts: Inputs, Loads, Process, Outputs, Never.

cyber5-command/

SKILL.md

config/tuned-values.md

stages/

00-setup · 01-ready · 02-plan

03-live — pace · sold-out · reroute · approval · scheduled

04-review

shared/ — engines + Meta tool notes

contracts/ — C5PLAN · C5LOG · handoff

output/ — templates · exports · deck.py

github.com/Fluency-Firm/Cyber5-Command

MIT licensed · samples included, anonymized